

## **Computer Networks and Distributed Systems**

### **Introduction**

Course 527 – Spring Term 2016-2017

**Emil C Lupu & Fidelis Perkonigg**

e.c.lupu@imperial.ac.uk, f.perkonigg@imperial.ac.uk

## **Course Structure**

1<sup>st</sup> half: **Distributed Systems** – covers basic distributed systems architectures, remote (object) interactions, remote procedure calls, security

2<sup>nd</sup> half: **Computer Networks** - covers basic principles of networking through examples of real technology

Whilst networks are concerned with communicating data from one endpoint to another (or several) distributed systems help us design systems whose components are hosted on different computers.

2

## **Course Structure**

2 lectures + 1 tutorial each week

1 coursework

Electronic handouts available from CATE

Please ask questions!

### **Acknowledgements:**

- Computer Networks based on material by Peter Pietzuch, Dan Chalmers and Ian Harries
- Distributed Systems based on material by Morris Sloman

## **Course Attendance**

B.Eng & M.Eng Electronic and Information Engineering  
2nd year      Required

B.Eng & M.Eng Mathematics and Computer Science  
3<sup>rd</sup> Year      Selective

M.Sc Computing Science      Selective

# Exam Questions

Not about low-level details

- Q: "What's the 50<sup>th</sup> bit in the IP packet header?"
- A: "It's the 'Don't Fragment' Flag"

But rather about principles and design trade-offs

- Q: "You need to design a transport layer for a network with the following characteristics... How would you do this?"
- A: "I'd use a reliable transport service, similar to TCP, because..."

Always explain your reasoning!

Some (simple) maths involved

5

# Recommended Books (for DS)

"Distributed Systems: Concepts and Design", George Coulouris, Jean Dollimore, Tim Kindberg, Gordon Blair Addison-Wesley

"Distributed Systems: Principles and Paradigms"  
Andrew S Tanenbaum, Maarten Van Steen, Pearson Ed.  
(2<sup>nd</sup> Ed.)

Acknowledgements:

- Distributed Systems based on material by Morris Sloman
- Some of the figures and images from external sources.

Imperial College  
London

## Computer Networks and Distributed Systems

### Part 2.1 – Distributed Systems Architecture

Course 527 – Spring Term 2016-2017

Emil C Lupu

e.c.lupu@imperial.ac.uk

[goo.gl/Dgesnp](http://goo.gl/Dgesnp)

Lecture notes

Tutorials + Notes on Solutions

Additional Material

Coursework in CATE

7

8

## Contents

Characteristics of Distributed Systems

Distributed Design

Peer-To-Peer

Layering

Transparencies

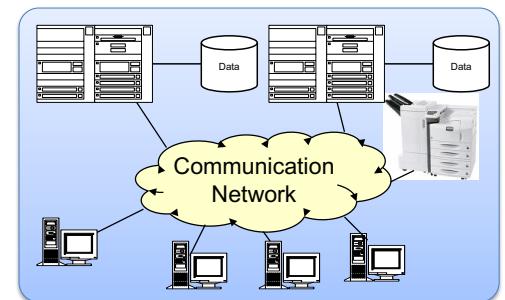
Viewpoints

## Definition

A distributed system consists of a collection of autonomous computers interconnected by a computer network and equipped with distributed system software to form an integrated computing facility.

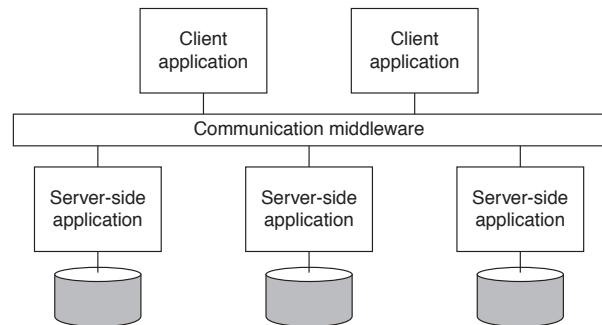
Components interact and cooperate to achieve a common goal.

Processes co-ordinate by means of **messages** transferred over a communication network.



9

## Enterprise Systems



Remote Procedure Calls

Message Oriented Middleware

Publish/Subscribe Systems

## Peer-to-peer (P2P)

Very large scale – potentially millions of users

Share processing eg Seti@home, United Devices, Avaki, Akamai

'Share' files eg Gnutella, Kazaa

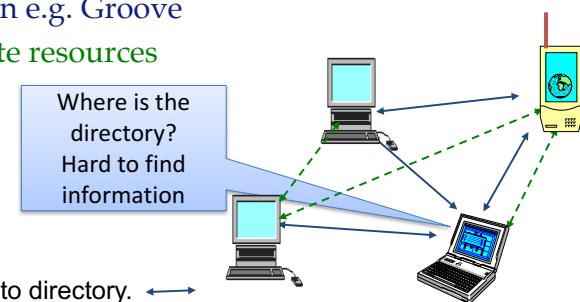
Collaboration e.g. Groove

How to locate resources

Where is the directory?  
Hard to find information

Publish and query to directory.

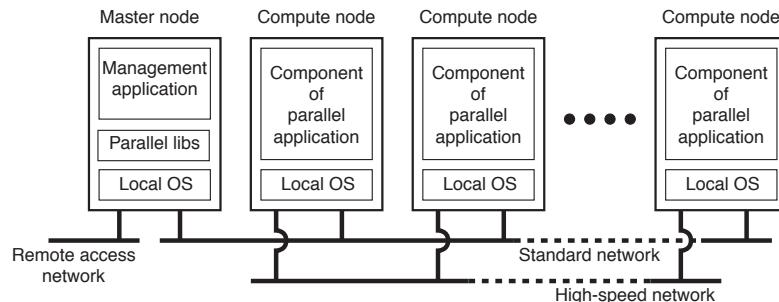
Get data from peers



11

12

## Computing Clusters



Generally used for parallel programming.  
Off-the-shelf computers interconnected by a high speed network.

Examples: Beowulf, CoW

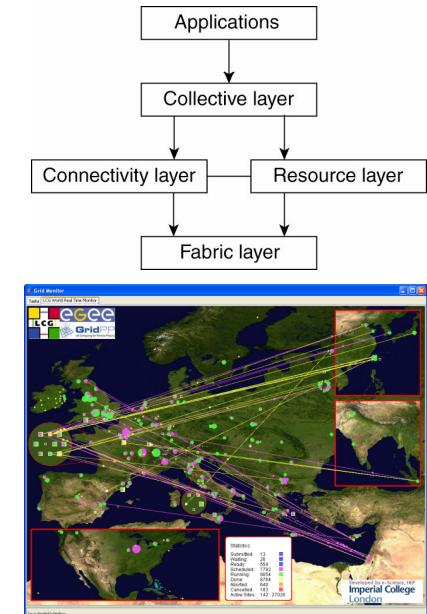


## Grid Computing

Distributing computation and data across several sites.

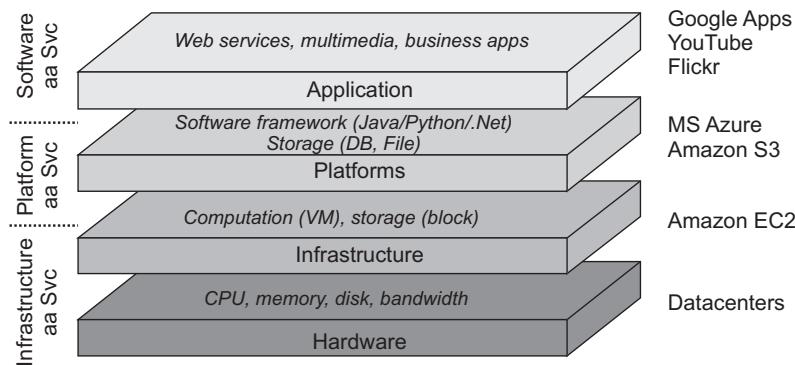
Increasingly based on Web Service Architectures

Examples: OGSA, Globus, OntoGrid



14

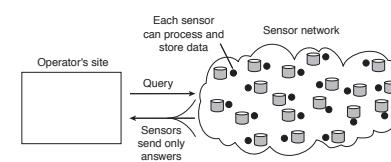
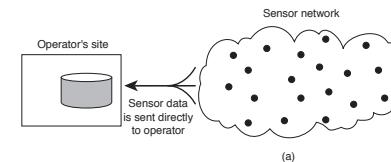
## Cloud Computing



Virtualisation of HW and SW infrastructure.

## (Distributed) Pervasive Systems IoT

### Wireless Sensor Networks

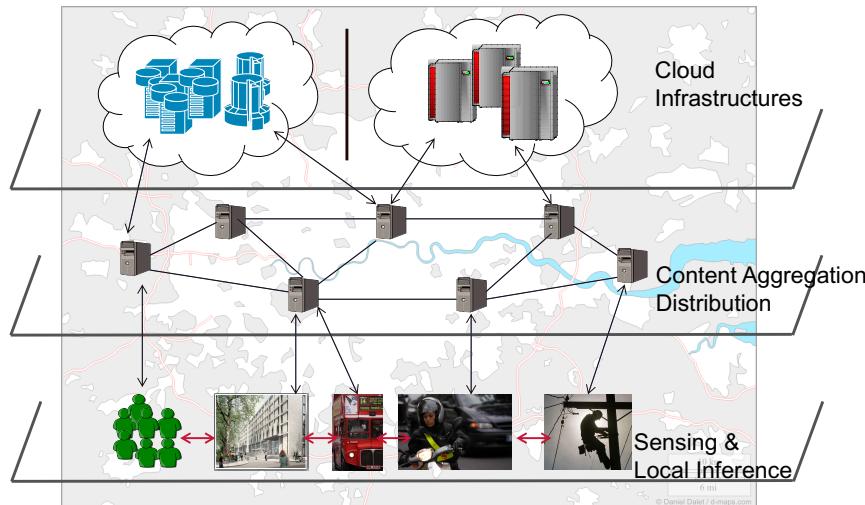


### Body Sensor Networks



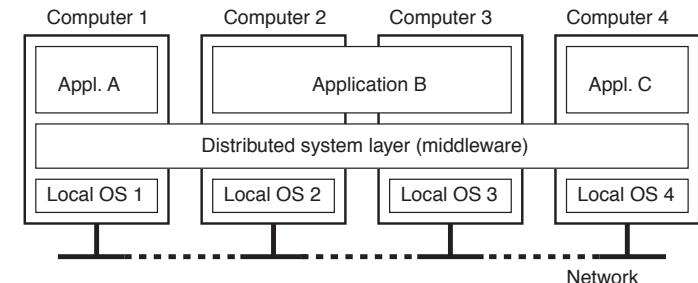
16

## Large Scale Infrastructures



17

## So What is a Distributed System?



Abstracting from the specificities of every node to consider the computing environment as a single coherent system.

Requires: interaction abstractions, middleware and services

18

## Characteristics/Advantages

**Resource sharing:** remote access to shared facilities

**Fault tolerance:** replication can remove failures

**Concurrency:** reduce response time by local processing, improve throughput by parallelism

**Openness:** vendor independence via clearly defined interfaces and use of standards

**Scalability:** via multiple processors and multiple networks

**Modularity:** simpler design, installation & maintenance

**Flexibility:** incremental change of function & adaptation to new requirements

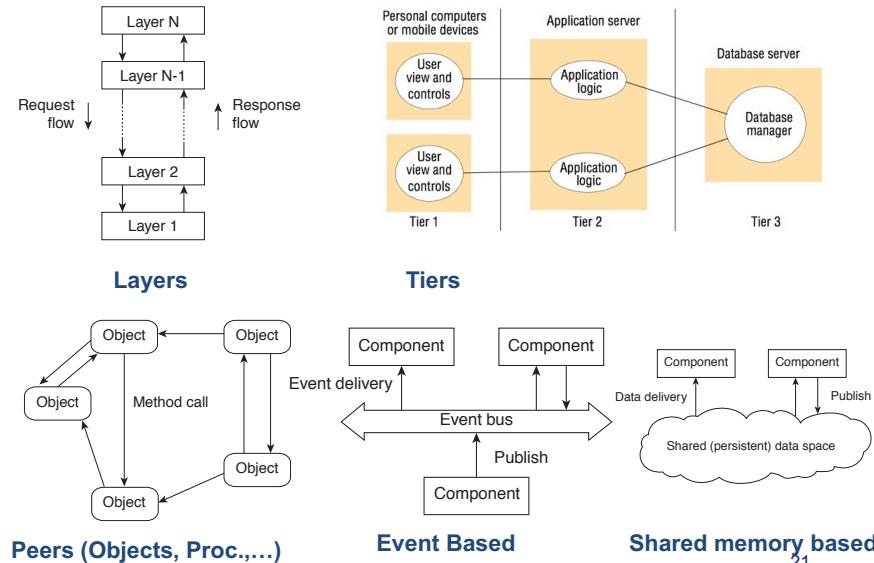
Reflect application distribution

**But no global time:** difficult to support causality and consistency

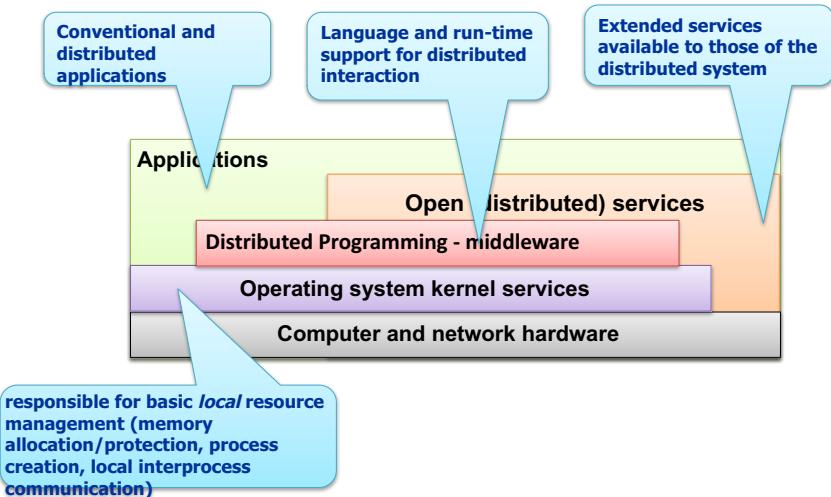
19

20

# Architectural Styles



# Basic software structure



22

## Open services:

- support the introduction of new services
- provide access to distributed services, including the coordination required for remote resource use (sharing, protection, synchronisation, recovery....)
- e.g. Jini resource discovery,

## Distributed programming support:

- supports interaction (such as remote procedure call) for conventional languages and support for special purpose languages.
- e.g. Java RMI, RPC

# Distribution Transparencies

It is often useful to hide distribution from the user – design goals which can be difficult to achieve

Access	uniform access whether local or remote
Location	access without knowledge of location ie location independent naming
Concurrency	sharing without interference – requires synchronisation
Replication	hides use of multiple components eg for fault tolerance
Failure	concealment of faults by replication or recovery
Migration	hides movement of components eg for load balancing
Performance	use of scheduling and reconfiguration to hide performance variations.
Scaling	permits expansion without changing system structure or algorithms
Heterogeneity	transforming information between different representations

23

24

# Open Distributed Processing –

## Viewpoints

Viewpoint = abstract representation of a system NOT phases in lifecycle model

### Enterprise Viewpoint

- Overall goals, policies & organisational structure
- Roles & activities within organisation(s)
- Policies & constraints regarding cross-organisation interactions
- Community: configuration of objects established to meet an objective – specifies roles, relationships and policies

### Information Viewpoint

- Modelling of information structures, information flows and knowledge representation
- Includes constraints on data
- No distinction between manual & automated information processing

25

### Computational Viewpoint

- Programming functions – IPC, object interfaces
- Application program structuring – independent of computer system on which it will run
- No distinction between processing & storage object
- Includes configuration – object instantiation and bindings.

### Engineering Viewpoint

- OS, communication system, database – implementation issues
- Provision of transparency mechanisms – fault tolerance, persistence etc.
- Processors & networks are visible

26

## Example: Pump for Mine Drainage

### Technology Viewpoint

- Realised components from which distributed systems are built.
- Particular OS (Unix, Windows), protocols (FTP, TCP/IP), processors (Intel, sparc, ARM)

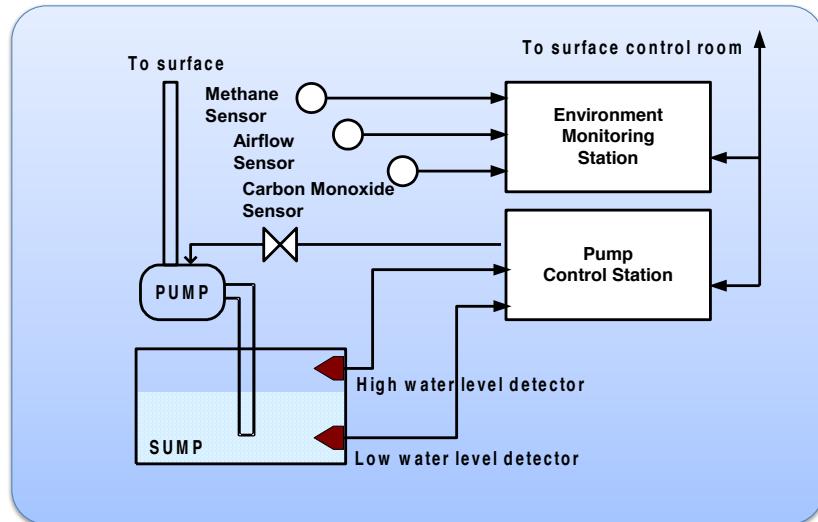
The pump is situated underground in a coal mine, and so for safety reasons it must not be started or continue running when the percentage of methane in the atmosphere exceeds a set safety limit. The pump controller obtains information on methane levels by communicating with a nearby environment monitoring station. As well as methane, this station also monitors carbon monoxide and airflow velocity. The environment monitoring station provides information to the surface and other plant controllers as well as to the pump controller.

Once start has been enabled by a command from the surface, the pump runs automatically controlled by the water level as sensed by the high and low level detectors. Detection of high level causes the pump to run until low level is reached. The surface may deactivate the pump with a stop command, and also query the status of the pump.

27

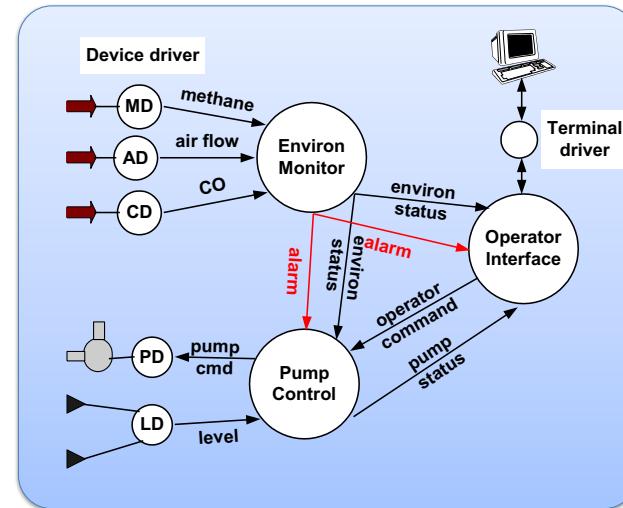
28

## Pump System Overview



29

## Pump Control Schematic



30

## Data Dictionary for Pump System

pump cmd	= (on, off)
level	= (high, low)
methane	= real
airflow	= real
alarm	= signal

environ status	= methane + airflow + CO
operator cmd	= (start, stop, status)
pump status	= (stopped, lowstop, methanestop, running)

31

### Data Flow Diagrams

Component Processes describe the functions of the system

Data flows = data type + direction

## Distributed System Design Approach

### Enterprise View

- Specify requirements and identify interactions with the environment
- Identify main processing components – processes or threads of control
- Assume 1 process per device

### Information View

- Identify data flows – direction and data types (dictionary)
- Ignore how interactions are initiated or types of interaction primitives

32

# Architecture Summary

## Computation View

- Decide on interaction primitives
- Decide on control flow i.e. whether data is pushed or pulled
- e.g. whether controllers are polled or event driven
- Specify component interfaces
- = interactions + signatures (parameter types)
- Specify component functions in terms of outline code and data structures

## Engineering View

- Optimise and allocate to physical nodes

## What are distributed systems

- definition
- potential benefits
- applications

## Architecture

- Viewpoint decomposition
- Architectural Style

## Main Design aspects